

```
In [1]: #Libraries
import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, KFold, cross_val_score, cv
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score, root_mean_squared_error
```

```
In [2]: #Reading the data
url = "https://raw.githubusercontent.com/angryolli/ML_Project/main/NBA.csv"
df = pd.read_csv(url, sep=',', on_bad_lines='skip')
df.head(1)
```

```
Out[2]:
```

	Player	Salary	Year	Pos	Age	Team	G	GS	MP	FG	...	FT%	ORB	DRB
0	Kobe Bryant	23034375	2010	SG	31	LAL	73	73	38.8	9.8	...	0.811	1.1	4.3

1 rows × 15 columns



```
In [3]: #Filtering for the last three seasons
#Couple of nan are in the dataset
df = df.query("Year >= 2023")
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1449 entries, 5849 to 7297
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype  
--- 
 0   Player    1449 non-null   object  
 1   Salary    1449 non-null   int64  
 2   Year      1449 non-null   int64  
 3   Pos       1449 non-null   object  
 4   Age       1449 non-null   int64  
 5   Team      1449 non-null   object  
 6   G         1449 non-null   int64  
 7   GS        1449 non-null   int64  
 8   MP        1449 non-null   float64 
 9   FG        1449 non-null   float64 
 10  FGA       1449 non-null   float64 
 11  FG%       1447 non-null   float64 
 12  3P        1449 non-null   float64 
 13  3PA       1449 non-null   float64 
 14  3P%       1429 non-null   float64 
 15  2P        1449 non-null   float64 
 16  2PA       1449 non-null   float64 
 17  2P%       1445 non-null   float64 
 18  eFG%      1447 non-null   float64 
 19  FT         1449 non-null   float64 
 20  FTA       1449 non-null   float64 
 21  FT%       1440 non-null   float64 
 22  ORB       1449 non-null   float64 
 23  DRB       1449 non-null   float64 
 24  TRB       1449 non-null   float64 
 25  AST        1449 non-null   float64 
 26  STL        1449 non-null   float64 
 27  BLK        1449 non-null   float64 
 28  TOV        1449 non-null   float64 
 29  PF         1449 non-null   float64 
 30  PTS        1449 non-null   float64 
dtypes: float64(23), int64(5), object(3)
memory usage: 362.2+ KB
```

```
In [5]: #columns including nan values
#delete players without throws
arr = ["FG%", "3P%", "2P%", "eFG%", "FT%"]
for a in arr:
    print("Null values " + a)
    print(df.loc[(df[a].isna()), ["Player", "FG", "FGA", "2P", "2PA", "3P", "3PA",
    df[[a]] = df[[a]].fillna(0)
df = df.drop([7267, 7275], axis=0)
```

Null values FG%

	Player	FG	FGA	2P	2PA	3P	3PA	FT	FTA
7267	Jalen McDaniels	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7275	Jahlil Okafor	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Null values 3P%

	Player	FG	FGA	2P	2PA	3P	3PA	FT	FTA
6824	Rudy Gobert	4.7	7.1	4.7	7.1	0.0	0.0	2.6	3.8
6920	Mitchell Robinson	2.2	3.3	2.2	3.3	0.0	0.0	0.8	1.1
6925	Daniel Gafford	5.0	7.1	5.0	7.1	0.0	0.0	2.4	3.4
6948	Ivica Zubac	7.4	11.8	7.4	11.8	0.0	0.0	2.0	3.0
7048	Dereck Lively II	3.7	5.3	3.7	5.3	0.0	0.0	1.3	2.0
7059	Jalen Duren	4.8	7.0	4.8	7.0	0.0	0.0	2.1	3.1
7063	Isaiah Jackson	2.8	4.6	2.8	4.6	0.0	0.0	1.4	2.8
7150	Jericho Sims	0.8	1.3	0.8	1.3	0.0	0.0	0.2	0.2
7155	Charles Bassey	1.9	3.3	1.9	3.3	0.0	0.0	0.6	0.9
7177	DeAndre Jordan	1.6	2.5	1.6	2.5	0.0	0.0	0.5	1.1
7230	Bruno Fernando	1.5	2.9	1.5	2.9	0.0	0.0	0.4	0.5
7231	Ariel Hukporti	0.8	1.2	0.8	1.2	0.0	0.0	0.2	0.5
7239	Ben Simmons	2.3	4.4	2.3	4.4	0.0	0.0	0.5	0.6
7259	Bismack Biyombo	2.4	4.1	2.4	4.1	0.0	0.0	0.4	0.9
7267	Jalen McDaniels	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7269	Terry Taylor	0.0	0.3	0.0	0.3	0.0	0.0	0.0	0.0
7275	Jahlil Okafor	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7276	Moses Brown	2.6	3.8	2.6	3.8	0.0	0.0	0.6	0.8
7286	Yuri Collins	0.5	1.5	0.5	1.5	0.0	0.0	0.0	0.0
7296	Kylor Kelley	1.3	2.0	1.3	2.0	0.0	0.0	0.5	0.9

Null values 2P%

	Player	FG	FGA	2P	2PA	3P	3PA	FT	FTA
7232	PJ Dozier	0.2	0.3	0.0	0.0	0.2	0.3	0.1	0.7
7267	Jalen McDaniels	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7271	Daishen Nix	0.0	0.3	0.0	0.0	0.0	0.3	0.3	0.7
7275	Jahlil Okafor	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Null values eFG%

	Player	FG	FGA	2P	2PA	3P	3PA	FT	FTA
7267	Jalen McDaniels	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7275	Jahlil Okafor	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

Null values FT%

	Player	FG	FGA	2P	2PA	3P	3PA	FT	FTA
7169	Joe Ingles	0.3	1.2	0.2	0.4	0.2	0.8	0.0	0.0
7178	Vlatko Cancar	0.8	2.0	0.5	0.8	0.3	1.2	0.0	0.0
7266	P.J. Tucker	1.0	2.3	0.0	0.3	1.0	2.0	0.0	0.0
7267	Jalen McDaniels	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7269	Terry Taylor	0.0	0.3	0.0	0.3	0.0	0.0	0.0	0.0
7270	Skal Labissiere	0.5	1.0	0.3	0.8	0.3	0.3	0.0	0.0
7275	Jahlil Okafor	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7286	Yuri Collins	0.5	1.5	0.5	1.5	0.0	0.0	0.0	0.0
7289	Terence Davis	0.0	2.0	0.0	1.0	0.0	1.0	0.0	0.0


```
In [6]: #0 nan values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1447 entries, 5849 to 7297
Data columns (total 31 columns):
 #   Column   Non-Null Count   Dtype  
 --- 
 0   Player    1447 non-null    object  
 1   Salary    1447 non-null    int64  
 2   Year      1447 non-null    int64  
 3   Pos       1447 non-null    object  
 4   Age       1447 non-null    int64  
 5   Team      1447 non-null    object  
 6   G         1447 non-null    int64  
 7   GS        1447 non-null    int64  
 8   MP        1447 non-null    float64 
 9   FG        1447 non-null    float64 
 10  FGA       1447 non-null    float64 
 11  FG%       1447 non-null    float64 
 12  3P        1447 non-null    float64 
 13  3PA       1447 non-null    float64 
 14  3P%       1447 non-null    float64 
 15  2P        1447 non-null    float64 
 16  2PA       1447 non-null    float64 
 17  2P%       1447 non-null    float64 
 18  eFG%      1447 non-null    float64 
 19  FT         1447 non-null    float64 
 20  FTA        1447 non-null    float64 
 21  FT%        1447 non-null    float64 
 22  ORB        1447 non-null    float64 
 23  DRB        1447 non-null    float64 
 24  TRB        1447 non-null    float64 
 25  AST        1447 non-null    float64 
 26  STL        1447 non-null    float64 
 27  BLK        1447 non-null    float64 
 28  TOV        1447 non-null    float64 
 29  PF         1447 non-null    float64 
 30  PTS        1447 non-null    float64 
dtypes: float64(23), int64(5), object(3)
memory usage: 361.8+ KB
```

```
In [7]: #Salary distribution
plt.hist(df.Salary, bins=30, range=(0, 55_000_000))
plt.xlabel("Salary $")
plt.ylabel("Number of players")
plt.title("Salary Distribution")
```

```
Out[7]: Text(0.5, 1.0, 'Salary Distribution')
```


Salary Distribution



```
In [8]: #Drop below the rookie salary players  
df = df[df["Salary"] >= 1157153]
```

```
In [9]: #Salary distribution  
plt.hist(df.Salary, bins=30, range=(0, 55_000_000))  
plt.xlabel("Salary $")  
plt.ylabel("Number of players")  
plt.title("Salary Distribution")
```

```
Out[9]: Text(0.5, 1.0, 'Salary Distribution')
```




```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1191 entries, 5849 to 7274
Data columns (total 31 columns):
 #   Column  Non-Null Count Dtype  
--- 
 0   Player    1191 non-null   object  
 1   Salary    1191 non-null   int64  
 2   Year      1191 non-null   int64  
 3   Pos       1191 non-null   object  
 4   Age       1191 non-null   int64  
 5   Team      1191 non-null   object  
 6   G         1191 non-null   int64  
 7   GS        1191 non-null   int64  
 8   MP        1191 non-null   float64 
 9   FG        1191 non-null   float64 
 10  FGA       1191 non-null   float64 
 11  FG%       1191 non-null   float64 
 12  3P        1191 non-null   float64 
 13  3PA       1191 non-null   float64 
 14  3P%       1191 non-null   float64 
 15  2P        1191 non-null   float64 
 16  2PA       1191 non-null   float64 
 17  2P%       1191 non-null   float64 
 18  eFG%      1191 non-null   float64 
 19  FT         1191 non-null   float64 
 20  FTA        1191 non-null   float64 
 21  FT%        1191 non-null   float64 
 22  ORB        1191 non-null   float64 
 23  DRB        1191 non-null   float64 
 24  TRB        1191 non-null   float64 
 25  AST        1191 non-null   float64 
 26  STL        1191 non-null   float64 
 27  BLK        1191 non-null   float64 
 28  TOV        1191 non-null   float64 
 29  PF         1191 non-null   float64 
 30  PTS        1191 non-null   float64 
dtypes: float64(23), int64(5), object(3)
memory usage: 297.8+ KB
```

```
In [11]: #Pearson correlation matrix
heat_salary= df[['Year', 'Age', 'Salary', 'G', 'GS', 'MP','PTS','AST','TRB','ORB',
'FG', 'FGA', 'FG%', '2P', '2PA', '2P%', '3P', '3PA', '3P%', 'FT', 'FTA', 'FT%', 'eFG%', '']
plt.figure(figsize=(14, 10))
dfData = heat_salary.corr()
sns.heatmap(dfData, annot=True, annot_kws={"size":6})
```

Out[11]: <Axes: >



```
In [12]: playerdb = df[["Salary", "Age", "G", "GS", "MP", "PTS", "AST", "TRB", "ORB", "DRB", "FG", "FGA", "FG%", "2P", "2PA", "2P%", "3P", "3PA", "3P%", "FT", "FTA", "FT%", "eFG%", "STL", "BLK", "TOV", "PF"]]

X = playerdb.drop(columns=["Salary"])
y = playerdb["Salary"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=49
)
```

```
In [13]: feature_sets = {
    "F1_important": ["PTS", "3PA", "DRB"],

    "F2_core": ["PTS", "MP", "TOV", "AST", "TRB", "STL", "BLK", "Age"],

    "F3_scoring": ["PTS", "3PA", "DRB", "STL", "MP", "FG"],

    "F4_all": [
        'Age', 'G', 'GS', 'MP', 'PTS', 'AST', 'TRB', 'ORB', 'DRB',
        'FG', 'FGA', 'FG%', '2P', '2PA', '2P%', '3P', '3PA', '3P%',
        'FT', 'FTA', 'FT%', 'eFG%', 'STL', 'BLK', 'TOV', 'PF'
    ]
}
```

```
In [14]: linreg_results = []

cv = KFold(n_splits=5, shuffle=True, random_state=49)

for name, features in feature_sets.items():
    X_tr = X_train[features]
    linreg = LinearRegression()
```

```
scores = cross_validate(
    linreg,
    X_tr,
    y_train,
    cv=cv,
    scoring={
        "rmse": "neg_root_mean_squared_error",
        "r2": "r2"
    }
)

linreg_results.append({
    "Model": "Linear Regression",
    "Feature_Set": name,
    "RMSE": -scores["test_rmse"].mean(),
    "R2": scores["test_r2"].mean()
})
```

```
In [15]: ranfor_results = []

for name, features in feature_sets.items():
    X_tr = X_train[features]

    rf = RandomForestRegressor(
        n_estimators=300,
        max_depth=None,
        random_state=49,
        n_jobs=-1
    )

    scores = cross_validate(
        rf,
        X_tr,
        y_train,
        cv=cv,
        scoring={
            "rmse": "neg_root_mean_squared_error",
            "r2": "r2"
        }
    )

    ranfor_results.append({
        "Model": "Random Forest",
        "Feature_Set": name,
        "RMSE": -scores["test_rmse"].mean(),
        "R2": scores["test_r2"].mean()
    })

ranfor_results
```

```
Out[15]: [ {'Model': 'Random Forest',
   'Feature_Set': 'F1_important',
   'RMSE': np.float64(8752335.068202568),
   'R2': np.float64(0.4796757551175661)},
  {'Model': 'Random Forest',
   'Feature_Set': 'F2_core',
   'RMSE': np.float64(6430518.963598682),
   'R2': np.float64(0.7221287812470895)},
  {'Model': 'Random Forest',
   'Feature_Set': 'F3_scoring',
   'RMSE': np.float64(8596308.319644738),
   'R2': np.float64(0.5024185227541278)},
  {'Model': 'Random Forest',
   'Feature_Set': 'F4_all',
   'RMSE': np.float64(6200264.348918025),
   'R2': np.float64(0.7421415110614864)}]
```

```
In [16]: results_df = pd.DataFrame(linreg_results + ranfor_results)
results_df["RMSE_mil"] = results_df["RMSE"] / 1e6
results_df.sort_values(by="RMSE")
#Best feature subset is F4_all for both models
```

	Model	Feature_Set	RMSE	R2	RMSE_mil
7	Random Forest	F4_all	6.200264e+06	0.742142	6.200264
5	Random Forest	F2_core	6.430519e+06	0.722129	6.430519
3	Linear Regression	F4_all	7.342383e+06	0.635035	7.342383
1	Linear Regression	F2_core	7.434146e+06	0.626774	7.434146
0	Linear Regression	F1_important	8.398112e+06	0.524606	8.398112
2	Linear Regression	F3_scoring	8.426318e+06	0.520983	8.426318
6	Random Forest	F3_scoring	8.596308e+06	0.502419	8.596308
4	Random Forest	F1_important	8.752335e+06	0.479676	8.752335

```
In [17]: wide_rmse = results_df.pivot(index="Feature_Set", columns="Model", values="RMSE")
wide_r2 = results_df.pivot(index="Feature_Set", columns="Model", values="R2")
wide_rmse
```

	Model	Linear Regression	Random Forest
Feature_Set			
F1_important	8.398112	8.752335	
F2_core	7.434146	6.430519	
F3_scoring	8.426318	8.596308	
F4_all	7.342383	6.200264	

```
In [18]: data = {
    "Model": wide_rmse.index.tolist(),           # feature sets
    "LR_RMSE": wide_rmse["Linear Regression"].tolist(),      # change "LR" to
    "LR_R2": wide_r2["Linear Regression"].tolist(),
```

```

    "RF_RMSE": wide_rmse["Random Forest"].tolist(),           # change "RF" to actu
    "RF_R2":   wide_r2["Random Forest"].tolist(),
}

ddf = pd.DataFrame(data)
x = np.arange(len(ddf["Model"]))
width = 0.35

def plot_model(title, rmse_values, r2_values):
    fig, ax1 = plt.subplots(figsize=(8,5))

    # RMSE bars
    bars1 = ax1.bar(x - width/2, rmse_values, width, color='blue', label='RMSE')
    ax1.set_ylabel("RMSE (lower is better)", color='blue')
    ax1.tick_params(axis='y', labelcolor='blue')

    # R^2 bars
    ax2 = ax1.twinx()      #second y-axis, same x-axis
    bars2 = ax2.bar(x + width/2, r2_values, width, color='orange', label='R^2')
    ax2.set_ylabel("R^2 (higher is better)", color='orange')
    ax2.tick_params(axis='y', labelcolor='orange')

    # X-axis labels
    ax1.set_xticks(x)
    ax1.set_xticklabels(ddf["Model"], rotation=30, ha='right')
    ax1.set_xlabel("Feature Set")

    # Add value labels
    for bar in bars1:
        ax1.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.15,
                 f"{bar.get_height():.2f}", ha='center', color='blue', fontsize=10)
    for bar in bars2:
        ax2.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.02,
                 f"{bar.get_height():.2f}", ha='center', color='darkorange', fontweight='bold', fontsize=10)

    # Adjust Y-axis upper limits for padding
    ax1.set_ylim(0, max(rmse_values) * 1.15)
    ax2.set_ylim(0, max(r2_values) * 1.15)

    # Get legend handles/labels from both axes
    handles1, labels1 = ax1.get_legend_handles_labels()
    handles2, labels2 = ax2.get_legend_handles_labels()



    # Title at the very top
    fig.suptitle(title, y=0.98)

    # Legend centered just under the title
    fig.legend(handles1 + handles2, labels1 + labels2,
               loc='upper center', bbox_to_anchor=(0.5, 0.90),
               ncol=2, frameon=True, framealpha=0.9, edgecolor='gray')

    # Leave room at the top for title + legend
    fig.tight_layout(rect=[0, 0, 1, 0.80])
    plt.show()

    # Plot both charts
    plot_model("Linear Regression: RMSE and R^2 Across Feature Sets", ddf["LR_RMSE"],
               plot_model("Random Forest: RMSE and R^2 Across Feature Sets", ddf["RF_RMSE"], ddf

```

Linear Regression: RMSE and R² Across Feature SetsRandom Forest: RMSE and R² Across Feature Sets

```
In [19]: best_features = feature_sets["F4_all"]
X_tr = X_train[best_features]
X_te = X_test[best_features]

rf = RandomForestRegressor(
    n_estimators=300,
    max_depth=None,
    random_state=49,
    n_jobs=-1
)
```

```
In [20]: rf.fit(X_tr, y_train)
y_pred = rf.predict(X_te)

test_rmse = root_mean_squared_error(y_test, y_pred)
```

```
test_r2 = r2_score(y_test, y_pred)

print("FINAL TEST RMSE:", test_rmse)
print("FINAL TEST R2:", test_r2)
```

FINAL TEST RMSE: 5685640.858484306
FINAL TEST R2: 0.7392993303249291